Minimum+1 (s,t)-cuts and dual edge sensitivity oracle

Surender Baswana ✉ ⓘD
Department of Computer Science & Engineering, IIT Kanpur, Kanpur – 208016, India

Koustav Bhanja ✉ ⓘD
Department of Computer Science & Engineering, IIT Kanpur, Kanpur – 208016, India

Abhyuday Pandey ✉
Department of Computer Science & Engineering, IIT Kanpur, Kanpur – 208016, India

—— **Abstract** ——————————————————————————————————

Let G be a directed multi-graph on n vertices and m edges with a designated source vertex s and a designated sink vertex t. We study the (s,t)-cuts of capacity minimum+1 and as an important application of them, we give a solution to the dual edge sensitivity for (s,t)-mincuts – reporting the (s,t)-mincut upon failure or addition of any pair of edges.

Picard and Queyranne [Mathematical Programming Studies, 13(1):8-16, 1980] showed that there exists a directed acyclic graph (DAG) that compactly stores all minimum (s,t)-cuts of G. This structure also acts as an oracle for the single edge sensitivity of minimum (s,t)-cut. Dinitz and Nutov [STOC, pages 509-518, 1995] showed that there exists an $\mathcal{O}(n)$ size 2-level cactus model that stores all global cuts of capacity minimum+1. However, for minimum+1 (s,t)-cuts, no such compact structures exist till date. We present the following structural and algorithmic results on minimum+1 (s,t)-cuts.

1. There exists a pair of DAGs of size $\mathcal{O}(m)$ that compactly store all minimum+1 (s,t)-cuts of G. Each minimum+1 (s,t)-cut appears as a (s,t)-cut in one of the 2 DAGs and is 3-transversal – it intersects any path in the DAG at most thrice.

2. There exists an $\mathcal{O}(n^2)$ size data structure that, given a pair of vertices $\{u,v\}$ which are not separated by an (s,t)-mincut, can determine in $\mathcal{O}(1)$ time if there exists a minimum+1 (s,t)-cut, say (A,B), such that $\{s,u\} \in A$ and $\{v,t\} \in B$; the corresponding cut can be reported in $\mathcal{O}(|B|)$ time.

3. There exists an $\mathcal{O}(n^2)$ size data structure that solves the dual edge sensitivity problem for (s,t)-mincuts. It takes $\mathcal{O}(1)$ time to report the value of a resulting (s,t)-mincut (A,B) and $\mathcal{O}(|B|)$ time to report the cut.

4. For the data structure problems addressed in (2) and (3) above, we also provide a matching conditional lower bound. We establish a close relationship among three seemingly unrelated problems – all-pairs directed reachability problem, the dual edge sensitivity problem for (s,t)-mincuts, and 2×2 maximum flow. Assuming the directed reachability hypothesis, this leads to $\tilde{\Omega}(n^2)$ lower bounds on the space for the latter two problems.

1 Introduction

The concept of cuts of a graph is very fundamental in graph theory and has many algorithmic applications as well. There are mainly two types of cuts – global cuts and (s,t)-cuts. A set of edges whose removal disconnects a given undirected graph is called a *global* cut. Let $G = (V, E)$ be a directed multi-graph (E is a multiset) consisting of $n = |V|$ vertices and

$m = |E|$ edges with a designated source vertex s and a designated sink vertex t. An (s,t)-cut of G is defined as follows.

▶ **Definition 1** ((s,t)-cut). *For a subset $C \subset V$ with $s \in C$ and $t \notin C$, the set of outgoing edges of C is called an (s,t)-cut.*

It follows from Definition 1 that for each (s,t)-cut, there exists at least one subset $C \subset V$ that defines it. For the simplicity of exposition and without causing any confusion, henceforth we shall use C to denote the corresponding (s,t)-cut as well. An edge is said to *contribute* to an (s,t)-cut C if it is an outgoing edge of C. The set of edges that have one endpoint in C and another endpoint in $\overline{C}$ is known as the *edge-set* of C, denoted by $E(C)$.

A (s,t)-cut (likewise a global cut) consisting of least number of contributing edges is called an (s,t)-mincut (likewise a global mincut).

There has been extensive research in designing efficient algorithms for global mincuts ([15], [18], [19]) as well as (s,t)-mincuts ([10], [16]). In addition, elegant graph structures have been invented that compactly store and characterize these cuts – Cactus graph for all global mincuts given by Dinitz, Karzanov, and Lomonosov [5], and a directed acyclic graph (DAG) given by Picard and Queyranne [23] for all (s,t)-mincuts. These structures can also serve as an efficient data structure for single edge sensitivity problem – to report the (s,t)-mincut (or global mincut) after the insertion/failure of an edge.

It is very natural to ask if there exists any compact structure for cuts of value greater than the value of the minimum cuts. For global minimum+1 cuts, Dinitz and Nutov answered this question in affirmative. In a seminal work [7], they showed that there exists an $\mathcal{O}(n)$ size 2-level cactus model that stores all minimum+1 cuts; they also gave a characterization of these cuts. An incremental maintenance of this structure solves the problem of maintaining minimum+2 edge connected components for any value of minimum cuts under insertion of edges; generalizing the results of Galil and Italiano [12] and Dinitz [6].

However, for minimum+1 (s,t)-cuts, to the best of our knowledge, no such compact structure exists till date. Note that the approach taken by Picard and Queyranne [23] for (s,t)-mincuts does not seem extendable for minimum+1 (s,t)-cuts. This is because their structure is based on the residual network resulting from a maximum (s,t)-flow and thus crucially exploits the equivalence between maximum flow and minimum cut (Ford and Fulkerson [10]); unfortunately, for a given minimum+1 (s,t)-cut, there is no equivalent (s,t)-flow.

While a compact structure for minimum+1 (s,t)-cuts is of significant importance from a graph theoretic perspective, equally important is a compact data structure that can efficiently answer the following fundamental query for any given edge $(u,v) \in E$.

$Q(u,v)$:   report a minimum+1 (s,t)-cut C, if exists, such that $u \in C$, $v \in \overline{C}$.

Interestingly, the DAG structure of Picard and Queyranne [23] for (s,t)-mincuts can answer efficiently the question stated above in case of (s,t)-mincuts. For this purpose, it crucially exploits the property that the set of (s,t)-mincuts are closed under intersection and union operations. Unfortunately, this property no longer holds for minimum+1 (s,t)-cuts, thus making it quite nontrivial to design a data structure for answering query Q.

The study of minimum+1 (s,t)-cuts has an immediate application as well. As is evident from the following fact, their study is indispensable for efficiently solving the dual edge sensitivity problem for (s,t)-mincuts.

▶ **Fact 1.1.** *The failure of a given pair of edges will reduce the value of (s,t)-mincut if and only if there is a minimum (s,t)-cut containing any failed edge or a minimum+1 (s,t)-cut containing both the failed edges.*

Sensitivity data structure for a given graph problem is motivated by the fact that graphs in the real world are prone to failures of vertices/edges. These failures are transient in nature. So while the set of failed vertices/edges keep changing with time, at any stage of time, the number of failed vertices/edges remains quite small. Therefore, the aim is to have a compact data structure that can efficiently report the solution of the given problem for a given set of failed vertices/edges. In a similar manner, we would like to efficiently report the solution of the given problem for a given set of newly added vertices/edges. This may help us determine which newly added vertex/edge changes the solution most significantly. In the past, many elegant fault-tolerant structures have been designed for various classical problems, like single-source reachability [17], shortest paths [4], breadth-first search [21], (s,t)-mincuts [23], all-pairs mincuts [1] etc, which can handle the failure of a single vertex/edge. It is certainly interesting and important to handle more than a single failure/addition. In this endeavour, it is quite natural to first design data structures that can handle dual failures (or dual insertions). This either helps, or exposes the difficulty, in solving the problem in its generality. It has turned out that the data structures that handle dual failures are often more complex and require deeper insight into the problem than the data structures that handle only single failure. This is evident at least for the following problems – single-source reachability [3], breadth-first search [20], shortest paths [9] etc. For the case of (s,t)-mincuts, note that the 40 years old data structure of Picard and Queyranne [23] is the only known sensitivity data structure and it can handle only a single edge failure/addition. It occupies $\mathcal{O}(m)$ space and can report the value of the resulting (s,t)-mincut and the corresponding cut in $\mathcal{O}(1)$ and $\mathcal{O}(m)$ time, respectively. No nontrivial data structure exists for the dual-edge sensitivity of (s,t)-mincuts till date.

1.1 New results and their overview

Let λ be the value of the (s,t)-mincut in G. Henceforth, we use $(\lambda + k)$ (s,t)-cut to denote a minimum+k (s,t)-cut, $k \in \{0,1\}$. Using just the sub-modularity of (s,t)-cuts (Lemma 9) and the relation between any pair of (s,t)-mincuts, we first present an alternate DAG structure, denoted by $\mathcal{D}_\lambda$, that compactly stores and characterizes all (s,t)-mincuts as follows. An (s,t)-cut in G is a (s,t)-mincut if and only if it appears as a 1-transversal cut in $\mathcal{D}_\lambda$ – the edge-set of an (s,t)-cut intersects any path in $\mathcal{D}_\lambda$ at most once. It can be easily observed that $\mathcal{D}_\lambda$, upon reversal of its edges, is identical to the DAG of Picard and Queyranne [23]. The major advantage of the approach taken for designing this alternate DAG is that it shows a way to design compact structures for $(\lambda + 1)$ (s,t)-cuts using only the properties of (s,t)-cuts without exploiting the relation between cuts and flows. We now present an overview of our main results on the $(\lambda + 1)$ (s,t)-cuts.

The set of minimum cuts are closed under intersection and union. This property has played a crucial role in designing compact structure as well as characterization of these mincuts – cactus graph for global mincuts by Dinitz, Karzanov, Lomonosov [5], the skeleton structure for Steiner mincuts by Dinitz and Vainshtein [8], and DAG for (s,t)-mincuts described in this paper. However, it turns out that $(\lambda + 1)$ (s,t)-cuts are not closed under intersection as well as union. Therefore, in order to design compact structure for $(\lambda+1)(s,t)$-cuts and their characterization, we analyse the relation between a pair of $(\lambda+1)(s,t)$-cuts. On the basis of the capacity of the cuts resulting from the intersection and union, each pair of $(\lambda + 1)$ (s,t)-cuts is classified into exactly one of the three types – Type-1, Type-2, and Type-3. While designing compact structure for $(\lambda + 1)$ (s,t)-cuts, the presence of pairs of Type-1 cuts poses a challenge in characterizing them. Likewise the presence of pairs of Type-2 cuts poses a challenge in determining the (s,t)-mincut upon the failure/addition of any pair of edges.

A natural idea to conquer these challenges is to partition the set of $(\lambda + 1)$ (s,t)-cuts into a *small* number of sets so that there is no pair of cuts from Type-1 (likewise Type-2) in a set. It can be observed that any arbitrary partitioning may produce a *large* number of sets which may defeat our objective of designing compact structures.

We present a technique called *covering* of (s,t)-cuts using which we can build a pair of graphs that cover all (s,t)-cuts of G and each of them has no pair of cuts from Type-1. It also helps in carefully handling pairs of cuts from Type-2. However, the covering technique works only for a graph with at most two (s,t)-mincuts, and hence can not be applied directly. In order to tackle this problem we introduce the concept of $(\lambda + 1)$ (s,t)-class as follows.

▶ **Definition 2.** *A $(\lambda + 1)$ (s,t)-class is a maximal set of vertices $A \subset V$ such that any pair of vertices from A are not separated by any (s,t)-mincut.*

It can be observed that the set of vertices of G that are assigned to a vertex of $\mathcal{D}_\lambda$ corresponds to a $(\lambda + 1)$ (s,t)-class. We first form a partition of the set of all $(\lambda + 1)$ (s,t)-cuts (excluding a set of *degenerate* cuts) with respect to the $(\lambda + 1)$ (s,t)-classes. For each $(\lambda + 1)$ (s,t)-class $\mathcal{W}$, we construct a new graph $G(\mathcal{W})$ that preserves all $(\lambda + 1)$ (s,t)-cut of G that subdivides $\mathcal{W}$ and has at most two (s,t)-mincuts. We show that it is sufficient to work with $G(\mathcal{W})$ for each $\mathcal{W}$ to design compact structure for $(\lambda + 1)$ (s,t)-cuts as well as dual edge sensitivity data structure for (s,t)-mincuts.

1. A 2-level DAG structure for $(\lambda + 1)$ (s,t)-cuts: Along similar lines of $\mathcal{D}_\lambda$, we build a compact graph $\mathcal{D}_{\lambda+1}$ that stores all $(\lambda + 1)$ (s,t)-cuts of $G(\mathcal{W})$. In order to establish the 1-transversality property of (s,t)-mincuts in $\mathcal{D}_\lambda$, the acyclicity of $\mathcal{D}_\lambda$ played a crucial role. Therefore, at first sight, we would expect $\mathcal{D}_{\lambda+1}$ to be acyclic and a $(\lambda + 1)$ (s,t)-cut C in $G(\mathcal{W})$ to be a ℓ-transversal cut in $\mathcal{D}_{\lambda+1}$ – edge-set of C intersects any path in $\mathcal{D}_{\lambda+1}$ at most ℓ times, for some constant ℓ. We find that $\mathcal{D}_{\lambda+1}$ is not necessarily acyclic and, surprisingly enough, a $(\lambda + 1)$ (s,t)-cut may appear in $\mathcal{D}_{\lambda+1}$ as an $\Omega(n)$-transversal cut. The root cause of non-acyclicity is the presence of pairs of cuts from Type-1. In order to tackle pairs of cuts from Type-1, we exploit the fact that $G(\mathcal{W})$ has at most two (s,t)-mincuts. We use the covering technique to construct a pair of graphs $G(\mathcal{W})^I$ and $G(\mathcal{W})^U$ that are Type-1 free. Now applying the same approach that was applied to obtain $\mathcal{D}_\lambda$, we construct a pair of DAGs – $\mathcal{D}_{\lambda+1}$ for graph $G(\mathcal{W})^I$ and $\mathcal{D}_{\lambda+1}$ for graph $G(\mathcal{W})^U$. This pair of DAGs is capable of characterizing each $(\lambda + 1)$ (s,t)-cut in G that subdivides $\mathcal{W}$ as follows. A $(\lambda + 1)$ (s,t)-cut in $G(\mathcal{W})$ appears as a *3-transversal* cut in exactly one of the two DAGs. In this way we obtain an $\mathcal{O}(m)$ size structure for compactly storing and characterizing all $(\lambda + 1)$ (s,t)-cuts and it consists of two levels – (i) DAG $\mathcal{D}_\lambda$ and (ii) a pair of DAGs for each $(\lambda + 1)$ (s,t)-class associated with a vertex of $\mathcal{D}_\lambda$.

Now we attempt to answer query $Q(u, v)$ using our 2-level DAG structure. Note that each 1-transversal cut in $\mathcal{D}_\lambda$ is also a (s,t)-mincut in G. Hence for the set of (s,t)-mincuts the query can be answered efficiently using a topological ordering of $\mathcal{D}_\lambda$. However, a 3-transversal cut in the pair of DAGs needs not be a $(\lambda + 1)$ (s,t)-cut because of the existence of *certain* pairs of cuts from Type-2.

2. Data structure for $(\lambda + 1)$ (s,t)-cuts: For each vertex u, there exists a unique (s,t)-mincut C, called *nearest (s,t)-mincut*, that keeps u on the side of s such that $C \subseteq C'$ for every other (s,t)-mincut C' that keeps u on the side of s. We can use this nearest (s,t)-mincut of u to determine if there exists a (s,t)-mincut C such that $u \in C$ and $v \in \overline{C}$ for any pair of vertices $\{u, v\}$. Unfortunately, there are multiple nearest $(\lambda + 1)$ (s,t)-cuts that keep u on the side of s, and hence this approach fails.

Let $\mathcal{W}$ be the $(\lambda+1)$ (s,t)-class to which u belongs. Let C and C' be any pair of nearest $(\lambda+1)$ (s,t)-cuts of u. We show that C and C' do not *cross* in $\mathcal{W}$, that is, $\overline{C \cup C'} \cap \mathcal{W} = \emptyset$. Using this crucial insight, we are able to design an $\mathcal{O}(n^2)$ size data structure summarized in the following theorem.

▶ **Theorem 3.** *Let G be a directed multi-graph on n vertices and m edges with a designated source vertex s and a designated sink vertex t. There exists a data structure occupying $\mathcal{O}(n^2)$ space that can determine in $\mathcal{O}(k)$ time whether there exists a $(\lambda+1)$ (s,t)-cut C such that $u \in C$ and $v_1, \ldots, v_k \in \overline{C}$ for any given vertices $u, v_1, \ldots, v_k$ belonging to a $(\lambda+1)$ (s,t)-class. It can also report C in $\mathcal{O}(|\overline{C}|)$ time.*

3. An oracle for dual edge sensitivity for (s,t)**-mincuts:** We show that there is a data structure $\{\mathcal{F}, \mathcal{I}\}$ occupying $\mathcal{O}(n^2)$ space which is capable of answering dual edge sensitivity query for (s,t)-mincuts in $\mathcal{O}(1)$ time. The data structure $\mathcal{F}$ for handling dual edge failure query is obtained as follows.

When both failed edges are not belonging to the same $(\lambda+1)$ (s,t)-class, data structures for (s,t)-mincuts are sufficient to answer the query. The main challenge arises when endpoints of both failed edges are belonging to the same $(\lambda+1)$ (s,t)-class $\mathcal{W}$. Notice that the data structure for $(\lambda+1)$ (s,t)-cut from Theorem 3 can determine whether there is a $(\lambda+1)$ (s,t)-cut in which a single failed edge is contributing, but cannot answer if both edges are contributing to a single $(\lambda+1)$ (s,t)-cut. Suppose there is a $(\lambda+1)$ (s,t)-cut C^* of G in which both failed edges, say (x,y) and (x',y'), are contributing. Then the necessary condition is that y' must not belong to the nearest $(\lambda+1)$ (s,t)-cut from x to y, and y must not belong to the nearest $(\lambda+1)$ (s,t)-cut from x' to y'. Therefore, if necessary condition holds then both edges are contributing to the union of the two nearest $(\lambda+1)$ (s,t)-cuts. However, the union needs not necessarily be a $(\lambda+1)$ (s,t)-cut because of the existence of certain pairs of $(\lambda+1)$ (s,t)-cuts from Type-2.

We employ the covering technique to tackle pairs of cuts from Type-2. Unfortunately, covering does not necessarily eliminate all Type-2 pairs of cuts like the way it does in case of Type-1 pairs. In order to tackle the problem arising due to the prevailing Type-2 pairs, we exploit the insight into the structure of $G(\mathcal{W})^I$ and $G(\mathcal{W})^U$. The end result is that, in order to determine whether any given pair of edges are contributing to a single $(\lambda+1)$ (s,t)-cut, we only have to perform a couple of nearest $(\lambda+1)$ (s,t)-cuts queries on $G(\mathcal{W})^I$ and $G(\mathcal{W})^U$.

4. Lower bound for various mincut data structures: We establish a close relationship between two seemingly unrelated problems – all-pairs directed reachability problem and dual edge sensitivity problem for (s,t)-mincuts. The classical problem of all-pairs directed reachability is defined as follows – Given a directed graph G on n vertices and m edges, preprocess it to form a data structure which can efficiently report if any given vertex v is reachable from another given vertex u. The problem becomes interesting when the underlying graph is sparse, that is, $m = o(n^2)$. It is natural to ask whether there is any data structure for the all-pairs directed reachability that takes $o(n^2)$ space and $o(m)$ query time ? Goldstein et al. [14], following the seminal work of Patrascu [22], stated a conjecture that concisely conveys the belief.

▶ **Conjecture 4** (Directed Reachability Hypothesis [14]). *Any data structure for the all-pairs reachability in a directed graph must either use $\tilde{\Omega}(n^2)$ space or $\tilde{\Omega}(m)$ query time. ($\tilde{\Omega}(.)$ denotes complexity upto polylogarithmic factors)*

We provide an $\mathcal{O}(m)$ time reduction from any instance of the all-pairs directed reachability problem for a given graph on n vertices and m edges to an instance of the dual edge sensitivity data structure for (s,t)-mincuts on a graph with $O(n)$ vertices and $\mathcal{O}(m)$ edges. Thus, assuming Conjecture 4 holds, our $\mathcal{O}(n^2)$ size data structure for dual edge sensitivity for (s,t)-mincuts is indeed the optimal size data structure for achieving $\mathcal{O}(1)$ query time. As a byproduct of this reduction, we establish conditional lower bounds on 2×2 flow tree. In particular, we show that if conjecture 4 holds, any data structure that can report the value of $(\{s,u\},\{v,t\})$-mincut for s,u,v and t being any four vertices of the graph must either use $\tilde{\Omega}(n^2)$ space or $O(m)$ query time. An interesting implication of this result is a matching conditional lower bound for the data structure described in Theorem 3.

1.2   Related work

Benczur [2] gave an $\mathcal{O}(n^2)$ space geometric representation for all the global cuts of value within $\frac{6}{5}$ times the global minimum cut value. As an important application, this structure leads to the improvement in the time complexity of *splitting* algorithms [11].

Vazirani and Yannakakis [24] addressed the following fundamental question about (s,t)-cuts. – Is there a polynomial time algorithm to compute an (s,t)-cut of second minimum weight? They answered this question in the affirmative by showing that there is an algorithm which can compute a k^{th} minimum weight (s,t)-cut using only $\mathcal{O}(n^{2(k-1)})$ maximum flow computations. This algorithm gives an implicit structure for all (s,t)-cuts – a binary tree with ℓ leaves, that stores all (s,t)-cuts of G and the number of (s,t)-cuts is ℓ.

Another related work is on characterizing (s,t)-cuts using the polyhedron corresponding to the dual of linear programming (LP) formulation on maximum (s,t)-flow. Vazirani and Garg [13] showed that not all (s,t)-cuts can be characterized as the vertices of this polyhedron. Moreover, they modify the dual of the LP formulation by adding a polynomial number of constraints such that the corresponding polyhedron of the resulting LP formulation has all (s,t)-cuts as its vertices.

1.3   Organisation of the paper

Section 2 contains the basic preliminaries and the construction of a quotient graph for a set of $(\lambda + k)$ (s,t)-cuts. The technique of covering (s,t)-cuts is explained in Section 3. Section 4 contains the construction and properties of the alternate DAG structure for (s,t)-mincuts. The structure for $(\lambda + 1)$ (s,t)-cuts and their characterization is discussed in Section 5. The data structure for reporting $(\lambda + 1)$ (s,t)-cuts is constructed in Section 6. Section 7 contains the Oracle for dual-edge sensitivity of (s,t)-mincuts. Finally in Section 8 we give the lower bounds. Missing proofs of lemmas and theorems are provided in the full version.

2   Preliminaries

For any $X, Y \subseteq V$, the capacity of (X,Y) (denote by $c(X,Y)$) is the number of edges leaving X and entering Y; for brevity we use $c(X)$ to denote $c(X, \overline{X})$ where $\overline{X} = (V \setminus X)$. We say that an (s,t)-cut C *subdivides* $X \subseteq V$ if $C \cap X$ and $\overline{C} \cap X$ are non empty.

▶ **Definition 5** (Nearest $(\lambda + k)$ (s,t)-cut of a vertex)**.** *The set $A \subset V$ with $s \in A$ and $t \in \overline{A}$ is said to be the nearest $(\lambda + k)$ (s,t)-cut of a vertex u if $u \in A$ and there is no $(\lambda + k)$ (s,t)-cut $B \subset V$ such that $u \in B$ and $B \subset A$. The set of all nearest $(\lambda + k)$ (s,t)-cuts of u is denoted by $N_k(u)$.*

▶ **Definition 6** (Nearest $(\lambda + k)$ (s, t)-cut for a pair of vertices). *Let A be a nearest $(\lambda + k)$ (s, t)-cut of a vertex u. For each vertex $v \in \overline{A}$, A is said to be a nearest $(\lambda + k)$ (s, t)-cut from vertex $\{u\}$ to a vertex $\{v\}$. The set of all nearest $(\lambda + k)$ (s, t)-cuts from u to v is denoted by $N_k(u, v)$.*

When $|N_1(u)| = 1$ (likewise $|N_1(u, v)| = 1$), without causing any ambiguity we use $N_1(u)$ (likewise $N_1(u, v)$) to denote the corresponding cut as well.

▶ **Definition 7** (ℓ-transversal cut). *A (s, t)-cut in a directed graph H is said to be ℓ-transversal, $\ell \geq 1$, if any path in H intersects with the edge-set of the (s, t)-cut at most ℓ times.*

▶ **Definition 8** (transpose of a graph). *Let H be a directed multi-graph with designated source vertex s and designated sink vertex t. The transpose of graph H (denoted by $(H)^{\mathcal{T}}$) is obtained by reversing the orientation of each edge of H with role of s and t swapped.*

▶ **Lemma 9** (Submodularity of cuts). *For any $A, B \subseteq V$, $c(A) + c(B) \geq c(A \cap B) + c(A \cup B)$.*

▶ **Lemma 10.** *Let A and B be any two (s, t)-mincuts in G that are crossing, that is, $A \setminus B$ as well as $B \setminus A$ are non-empty. There is no edge of G between $A \setminus B$ and $B \setminus A$.*

2.1 Quotient graph for a family of (s, t)-cuts

Let $\mathcal{C}$ be a set of $(\lambda + k)$ (s, t)-cuts in graph $G = (V, E)$ where $k \in \{0, 1\}$. We define the following binary relation on the vertex set of G.

▶ **Definition 11** (Relation $R_{\lambda+k}$). *Any two vertices $\{x, y\} \in V$ are said to be related by $R_{\lambda+k}$ if and only if x and y are not separated by any $(\lambda + k)$ (s, t)-cut from $\mathcal{C}$.*

It is a simple exercise to show that $R_{\lambda+k}$ defines an equivalence relation on the vertex set. We call each equivalence class defined by $R_{\lambda+k}$ as a $(\lambda + k + 1)$ (s, t)-class. It can be observed that any (s, t)-cut that subdivides a $(\lambda + k + 1)$ (s, t)-class has capacity strictly larger than $\lambda + k$. The following lemma states how a $(\lambda + k)$ (s, t)-class is related to a $(\lambda + k)$ (s, t)-cut.

▶ **Lemma 12.** *A $(\lambda + k)$ (s, t)-cut can subdivide at most one $(\lambda + k)$ (s, t)-class.*

Since $R_{\lambda+k}$ is an equivalence relation, the $(\lambda + k + 1)$ (s, t)-classes form a partition of the vertex set of G into disjoint subsets. Let $G_{\lambda+k}$ be the quotient graph of G obtained by contracting each of these subsets into single vertices. We call each vertex of the quotient graph as nodes. The node of $G_{\lambda+k}$ containing source s is denoted by S and the node containing sink t is denoted by T. We call an (S, T)-cut of $G_{\lambda+k}$ as (s, t)-cut without causing any ambiguity. The following theorem is immediate from the construction.

▶ **Theorem 13.** *For a directed multi-graph G and a set of $(\lambda + k)$ (s, t)-cuts $\mathcal{C}$ for any $k \in \{0, 1\}$, there exists a quotient graph $G_{\lambda+k}$ of G such that $C \in \mathcal{C}$ is an (s, t)-cut in G if and only if C is a $(\lambda + k)$ (s, t)-cut in $G_{\lambda+k}$.*

An edge of $G_{\lambda+k}$ is classified as normal or *inverted* using the following definition.

▶ **Definition 14** (inverted edge for $(\lambda + k)$ (s, t)-cut). *An edge (x, y) of $G_{\lambda+k}$ is said to be an inverted edge if there exists no $(\lambda + k)$ (s, t)-cut $C \in \mathcal{C}$ such that $x \in C$ and $y \in \overline{C}$.*

▶ **Fact 2.1.** *For a directed multi-graph H and a set of $(\lambda + k)$ (s, t)-cuts $\mathcal{C}$, let $H_{\lambda+k}$ be the quotient graph of H. Given a pair of nodes μ, ν in graph $H_{\lambda+k}$, the contraction of all vertices which are mapped to μ or ν into a single vertex in H eliminates precisely all those $(\lambda + k)$ (s, t)-cuts from $\mathcal{C}$ that separate μ and ν.*

Consider two sets of $(\lambda + k)$ (s,t)-cuts $\mathcal{C}$ and $\mathcal{C}'$ such that $\mathcal{C} \subset \mathcal{C}'$. Let $G_{\lambda+k}$ and $G'_{\lambda+k}$ be the quotient graphs for $\mathcal{C}$ and $\mathcal{C}'$ respectively. It follows from construction of quotient graphs that there exists equivalence classes of $G_{\lambda+k}$ which are further split into multiple equivalence classes of $G'_{\lambda+k}$. Therefore, the following lemma is immediate.

▶ **Lemma 15.** *Given a directed multi-graph G, let $G_{\lambda+k}$ and $G'_{\lambda+k}$ be a pair of quotient graphs formed on the sets of $(\lambda + k)$ (s,t)-cuts $\mathcal{C}$ and $\mathcal{C}'$ respectively. If $\mathcal{C} \subset \mathcal{C}'$ then $G_{\lambda+k}$ is a quotient graph of $G'_{\lambda+k}$.*

3   A covering of all (s, t)-cuts of a special graph

We first formalize the notion of *covering* the (s,t)-cuts of a graph using the following definition.

▶ **Definition 16** (Covering all (s,t)-cuts). *Given a directed multi-graph H with a designated source vertex s and a designated sink vertex t, let $F = \{H^1, H^2, \ldots, H^\ell\}$ be a finite set of directed multi-graphs, each defined on the same vertex set as that of H. F is said to cover all (s,t)-cuts of H if and only if the following conditions hold.*

1. *For each (s,t)-cut C in H, there exists a (s,t)-cut C' in H^i for a unique $1 \le i \le \ell$ such that $C = C'$.*
2. *For each (s,t)-cut C of finite capacity in H^i for any $1 \le i \le \ell$, there exists a (s,t)-cut C' in H such that $C = C'$.*

Let H be a directed multi-graph with a designated source vertex s and a designated sink vertex t that has at most two (s,t)-mincuts – $\{s\}$ and (complement of) $\{t\}$. Our aim is to compute a small family $F = \{H^1, H^2, \ldots, H^\ell\}$ that covers all (s,t)-cuts of H and satisfies the following condition – for each pair C, C' of (s,t)-cuts in H^i for any $1 \le i \le \ell$, either their intersection or union is not a λ (s,t)-cut.

We shall show that there exists a family $F = \{H^U, H^I\}$ of only two graphs that *covers* all (s,t)-cuts of H and satisfies the said property. We build two graph H^U and H^I from H as follows. Let x be any arbitrary vertex other than s and t of H. H^I is formed by adding infinite capacity edge from vertex s to x. In a similar way, H^U is formed by adding infinite capacity edge from vertex x to t. The following lemma ensures that H^U and H^I together cover all $(\lambda + k)$ (s,t)-cuts of H. It exploits the following simple fact – Let x be a vertex of graph H and C be any (s,t)-cut; either $x \in C$ or $x \in \overline{C}$.

▶ **Lemma 17.** *Let $\lambda' > 0$ be a finite number. C is an (s,t)-cut in H of capacity λ' if and only if C is a cut of capacity λ' in H^U or H^I.*

It follows from the construction that the first (s,t)-mincut of H, that is $\{s\}$, is present in H^U while the second (s,t)-mincut of H, that is (complement of) $\{t\}$ is present in H^I. So using Lemma 17, we can state the following theorem.

▶ **Theorem 18.** *Let $H = (V, E)$ be a directed multi-graph on n vertices and m edges with a designated source vertex s and a designated sink vertex t. Suppose H has at most two (s,t)-mincuts, $\{s\}$ and $V \setminus \{t\}$. There exists a set of graphs $F = \{H^I, H^U\}$ satisfying the following properties.*

1. *F covers all $(\lambda + k)$ (s,t)-cuts of H.*
2. *$V \setminus \{t\}$ is the only λ (s,t)-cut in H^I, and $\{s\}$ is the only λ (s,t)-cut in H^U.*

▶ **Note 19.** Covering technique cannot be applied to a graph H that has more than two (s, t)-mincuts. This is because, for any vertex x in H, there will be more than one (s, t)-mincuts that keep x on side of s and/or more than one (s, t)-mincuts that keep x on side of t. Therefore, there will be multiple (s, t)-mincuts in H^I or multiple (s, t)-mincuts in H^U, hence violating Theorem 18(2).

4 An alternate DAG structure storing all (s, t)-mincuts

We now present an alternate compact structure $\mathcal{D}_\lambda$ for representing and characterizing all (s, t)-mincuts of G.

Construction of $\mathcal{D}_\lambda$: Let G_λ be the quotient graph defined by all λ (s, t)-cuts of G (Theorem 13). D_λ is the graph obtained by reversing the direction of each inverted edge of G_λ. Exploiting just the elementary properties of (s, t)-mincuts alone (Lemma 10 and submodularity of cuts), we show that D_λ is acyclic. As the reader may notice, D_λ is related to the DAG of Picard and Queyranne [23] – the mapping of vertices to the nodes of the DAG is identical while the direction of each edge is flipped. Each (s, t)-mincut of G is characterized by a 1-transversal cut in D_λ as stated in the following lemma.

▶ **Lemma 20** (1-transversality property). *An (s, t)-cut is an (s, t)-mincut of G if and only if it appears as a 1-transversal cut in $\mathcal{D}_\lambda$.*

The following lemma can be proved along similar lines as Lemma 20.

▶ **Lemma 21.** *An (s, t)-cut of G is an (s, t)-mincut if and only if it appears as an (s, t)-cut with no incoming edge in $\mathcal{D}_\lambda$.*

Interestingly, $\mathcal{D}_\lambda$ can serve as a compact data structure for efficiently answering the query $Q(u, v)$ for (s, t)-mincuts as follows. Let τ be a topological ordering of $\mathcal{D}_\lambda$. For each edge (μ, ν) of $\mathcal{D}_\lambda$, $\tau(\mu) < \tau(\nu)$. Hence, it is a simple corollary of Lemma 20 that every prefix of τ is a (s, t)-mincut. Therefore, given any edge $(u, v) \in E$, there is a (s, t)-mincut in G containing it if u and v appear in different nodes of $\mathcal{D}_\lambda$ and the node containing u appears before the node containing v.

5 Structure of $(\lambda + 1)$ (s, t)-cuts

In order to construct a compact structure for $(\lambda + 1)$ (s, t)-cuts, we take an approach similar to the construction of $\mathcal{D}_\lambda$ for (s, t)-mincuts and construct a graph $\mathcal{D}_{\lambda+1}$ as follows.

Construction of $\mathcal{D}_{\lambda+1}$: Let $G_{\lambda+1}$ be the quotient graph defined by all $(\lambda + 1)$ (s, t)-cuts (Theorem 13). $D_{\lambda+1}$ is the graph obtained by reversing the direction of each inverted edge of $G_{\lambda+1}$. Henceforth $\mathcal{D}_{\lambda+1}$ for a graph H is denoted by $\mathcal{D}_{\lambda+1}(H)$.

$\mathcal{D}_\lambda$ characterizes (s, t)-mincuts of G as 1-transversal cuts. Unfortunately, it turns out that $\mathcal{D}_{\lambda+1}(G)$ is not sufficient for characterizing $(\lambda + 1)$ (s, t)-cuts in terms of ℓ-tranversal cuts for some constant ℓ as stated in the following theorem.

▶ **Theorem 22** ($\Omega(n)$-transversality). *There exists a directed multi-graph H on n vertices with a designated source vertex s and a designated sink vertex t having only two (s, t)-mincuts and a $(\lambda + 1)$ (s, t)-cut C such that C appears as a $\Omega(n)$-transversal cut in $\mathcal{D}_{\lambda+1}(H)$.*

We now provide a classification of all-pairs of $(\lambda + 1)$ (s, t)-cuts. This classification will act as a crucial tool in the analysis of $\mathcal{D}_{\lambda+1}(G)$ as well as establishing properties of the compact structure for $(\lambda + 1)$ (s, t)-cuts in Section 5.1.

■ **Figure 1** Contributing edges of a pair of (s,t)-cuts A and B between different regions.

Suppose A and B are any two $(\lambda+1)$ (s,t)-cuts. Using sub-modularity of cuts (Lemma 9), we arrive at the following inequality, $c(A \cap B) + c(A \cup B) \le c(A) + c(B) = 2(\lambda+1)$. Note that $c(A \cap B) \ge \lambda$ and $c(A \cup B) \ge \lambda$. This implies that the value of $c(A \cap B) + c(A \cup B)$ will always belong to $\{2\lambda, 2\lambda+1, 2\lambda+2\}$. Therefore, each pair (A, B) of $(\lambda+1)$ (s,t)-cuts can be classified uniquely into one of the three types based on the value of $c(A \cap B) + c(A \cup B)$. We now state the following lemma that will provide an alternate characterization of these three types based on the number of edges between $A \setminus B$ and $B \setminus A$.

▶ **Lemma 23.** *Suppose A and B are any two $(\lambda+1)$ (s,t)-cuts. Let γ_1 and γ_2 denote the number of edges from $A \setminus B$ to $B \setminus A$ and from $B \setminus A$ to $A \setminus B$ respectively. If $c(A \cap B) = \lambda + p$ and $c(A \cup B) = \lambda + q$ for some $p, q \ge 0$, then $\gamma_1 + \gamma_2 = 2 - p - q$.*

Proof. The cuts A and B partition G into at most 4 regions. Refer to Figure 1 for the illustration of these regions and the edges contributing to the respective cuts.

The following equations follow directly from Figure 1.

$$c(A \cap B) = \alpha_1 + \alpha_2 + \delta = \lambda + p \tag{1}$$
$$c(A \cup B) = \beta_1 + \beta_2 + \delta = \lambda + q \tag{2}$$
$$c(B) = \beta_1 + \alpha_2 + \gamma_2 + \delta = \lambda + 1 \tag{3}$$
$$c(A) = \beta_2 + \alpha_1 + \gamma_1 + \delta = \lambda + 1 \tag{4}$$

Using equations (1) and (3) we get $\alpha_1 = \beta_1 + \gamma_2 + p - 1$. Similarly using equations (1) and (4) we get $\alpha_2 = \beta_2 + \gamma_1 + p - 1$. By combining these two equalities, we get $\alpha_1 + \alpha_2 = \beta_1 + \beta_2 + \gamma_1 + \gamma_2 + 2p - 2$. Hence $\gamma_1 + \gamma_2 = 2 - p - q$, since $\alpha_1 + \alpha_2 = \beta_1 + \beta_2 + p - q$ from equations (2) and (1). ◀

Refer to Table 1 for two ways to characterize the three types of pairs of $(\lambda+1)$ (s,t)-cuts.

	$c(A \cap B) + c(A \cup B)$	$c(A \setminus B, B \setminus A) + c(B \setminus A, A \setminus B)$
Type-1	2λ	2
Type-2	$2\lambda + 2$	0
Type-3	$2\lambda + 1$	1

■ **Table 1** Classification of any pair (A, B) of $(\lambda+1)$ (s,t)-cuts.

Figure 2 (i) An example of $\mathcal{D}_{\lambda+1}(G)$ with cycle $\langle a, b, c, d, e, f, a \rangle$. Note that here $\mathcal{D}_{\lambda+1}(G)$ turns out to be same as G. (ii) Cycle O in graph $\mathcal{D}_{\lambda+1}(H)$ from the proof of Lemma 25.

It follows from Theorem 22 that transversality of $(\lambda+1)$ (s,t)-cuts cannot be bounded by a constant even if a graph has at most two (s,t)-mincuts – $\{s\}$ and (complement of) $\{t\}$. In the following section, we shall first give a compact structure for $(\lambda + 1)$ (s,t)-cuts for graph G assuming that G has at most two (s,t)-mincuts. Finally, we shall extend it to general graphs with the help of the structure of $\mathcal{D}_\lambda$ in Section 5.1.1.

5.1 Compact representation and characterization of $(\lambda + 1)$ (s,t)-cuts

The characterization of (s,t)-mincuts crucially exploits the fact that $\mathcal{D}_\lambda$ is a DAG. Unfortunately, as shown in Figure 2(i), $\mathcal{D}_{\lambda+1}(G)$ can have cycles. The following lemma states the source of any 2-length cycle in $\mathcal{D}_{\lambda+1}$.

▶ **Lemma 24.** *If graph G does not have any pair of cuts from Type-1, then $\mathcal{D}_{\lambda+1}(G)$ cannot have a cycle of length two.*

Proof. Suppose we have a 2-length cycle $\langle \mu, \nu, \mu \rangle$ in $\mathcal{D}_{\lambda+1}(G)$. It follows from construction of $\mathcal{D}_{\lambda+1}$ that there exists a $(\lambda + 1)$ (s,t)-cut C for the edge (μ, ν) such that $\mu \in C$ and $\nu \in \overline{C}$. In a similar manner there exists a $(\lambda + 1)$ (s,t)-cut C' for the edge (ν, μ) such that $\nu \in C'$ and $\mu \in \overline{C'}$. It can be observed that (C, C') forms a pair of cuts from Type-1, a contradiction. ◀

In fact, the absence of pairs of $(\lambda + 1)$ (s,t)-cuts from Type-1 is a sufficient condition for acyclicity as stated in the following lemma.

▶ **Lemma 25** (acyclicity property). *If graph G does not have any pair of cuts from Type-1, then $\mathcal{D}_{\lambda+1}(G)$ is a directed acyclic graph.*

Proof. We begin with stating the following assertion.
$\mathcal{A}(i)$: For any graph H that has no pair of $(\lambda + 1)$ (s,t)-cuts from Type 1, there is no cycle of length i in $\mathcal{D}_{\lambda+1}(H)$.
We shall now prove, by induction on i, that $\mathcal{A}(i)$ holds for all $i \geq 2$, and this would establish the lemma. The base case, $\mathcal{A}(2)$ follows directly from Lemma 24. Suppose $\mathcal{A}(j)$ holds for all $j < i$. We shall now prove that $\mathcal{A}(i)$ holds.

Suppose there is a cycle O of length i in $\mathcal{D}_{\lambda+1}(H)$ (see Figure 2(ii)). Consider any arbitrary edge (μ, ν) in this cycle. It follows from construction of $\mathcal{D}_{\lambda+1}(H)$ that there exists a $(\lambda + 1)$ (s,t)-cut C of H such that $\mu \in C$ and $\nu \in \overline{C}$. Let $\mathcal{U}$ be the set of vertices of H that are mapped to either μ or ν. Contracting the set $\mathcal{U}$ into a single vertex we obtain a new graph H' from H.

It follows from Fact 2.1 and Lemma 15 that $\mathcal{D}_{\lambda+1}(H')$ has to be a quotient graph of $\mathcal{D}_{\lambda+1}(H)$. As a result, cycle O in $\mathcal{D}_{\lambda+1}(H)$ will be mapped to either a cycle of length strictly less than i or a single node in $\mathcal{D}_{\lambda+1}(H')$. The Induction Hypothesis rules out the possibility of the former case. We shall now rule out the possibility of the latter case. This will establish the validity of $\mathcal{A}(i)$.

The cut C must contain at least one more edge, say (μ', ν'), of cycle O because the cycle O must intersect the edge-set of C at least twice. It follows from the construction of $\mathcal{D}_{\lambda+1}(H)$ that there exists a $(\lambda+1)$ (s,t)-cut, say C', in $\mathcal{D}_{\lambda+1}(H)$ such that $\mu' \in C'$ and $\nu' \in \overline{C'}$. Since μ' and ν' belong to the same node in $\mathcal{D}_{\lambda+1}(H')$, the cut C' is not present in H'. This observation, in the light of Fact 2.1, implies that C' must separate μ and ν (dashed curve in Figure 2). This would imply that for the pair of cuts (C, C'), $c(C \setminus C', C' \setminus C) + c(C' \setminus C, C \setminus C') = 2$. Hence (C, C') forms a Type-1 pair of $(\lambda+1)$ (s,t)-cut in H, a contradiction. ◄

In order to have acyclic structures that collectively preserve all $(\lambda+1)$ (s,t)-cuts, Lemma 25 raises the following question – can we partition the set of $(\lambda+1)$ (s,t)-cuts such that each partition does not contain any pair of $(\lambda+1)$ (s,t)-cut from Type-1? The covering technique (Theorem 18) gives an affirmative answer to this question. In particular, it outputs just a pair of graphs $\{G^I, G^U\}$ such that all $(\lambda+1)$ (s,t)-cuts of G are covered by G^I and G^U together. Moreover, both G^U and G^I will contain exactly one (s,t)-mincut each, and this ensures that there are no pairs of $(\lambda+1)$ (s,t)-cuts from Type-1 in G^I or G^U. Therefore, it follows from Lemma 25 that $\mathcal{D}_{\lambda+1}(G^I)$ and $\mathcal{D}_{\lambda+1}(G^U)$ are acyclic. In the case when G has exactly one (s,t)-mincut, $\mathcal{D}_{\lambda+1}(G)$ itself is acyclic.

Not only $\mathcal{D}_{\lambda+1}(G^I)$ and $\mathcal{D}_{\lambda+1}(G^U)$ are acyclic but also help in characterizing $(\lambda+1)$ (s,t)-cuts as follows.

We show that each $(\lambda+1)$ (s,t)-cut of G is 3-transversal in $\mathcal{D}_{\lambda+1}(G^I)$ or in $\mathcal{D}_{\lambda+1}(G^U)$. Without loss of generality let us consider the graph $\mathcal{D}_{\lambda+1}(G^U)$. In order to accomplish 3-transversality of $(\lambda+1)$ (s,t)-cuts in $\mathcal{D}_{\lambda+1}(G^U)$, we first show that for any path in $\mathcal{D}_{\lambda+1}(G^U)$ whose first node is not S, there is no $(\lambda+1)$ (s,t)-cut which can keep both the first and the last node of the path on side of S, and remaining nodes on side of T. As a warm-up, the following lemma validates this assertion for all paths of length two.

▶ **Lemma 26.** *Let $P = \langle v_0 \neq S, v_1, v_2 \rangle$ be a path in $\mathcal{D}_{\lambda+1}(G^U)$. There can not exist any $(\lambda+1)$ (s,t)-cut C such that $v_0, v_2 \in C$ and $v_1 \in \overline{C}$.*

Proof. We give a proof by contradiction. Assume that there is such a $(\lambda+1)$ (s,t)-cut C. It follows from the construction of $\mathcal{D}_{\lambda+1}(G^U)$ that there is a $(\lambda+1)$ (s,t)-cut C' for the edge (v_1, v_2) such that $v_1 \in C'$ and $v_2 \in \overline{C'}$. It can be observed that the edge (v_1, v_2) is an edge from $C' \setminus C$ to $C \setminus C'$. Therefore, $\{C, C'\}$ cannot be a pair from Type-2. $\{C, C'\}$ cannot be pair from Type-1 as well since there is no pair of $(\lambda+1)$ (s,t)-cuts from Type-1 in $\mathcal{D}_{\lambda+1}(G^U)$. So, $\{C, C'\}$ must be a pair from Type-3. Notice that $C \cup C'$ has capacity $\lambda+1$ since the only λ (s,t)-cut in G^U is $\{s\}$. Hence, $C \cap C'$ must be containing the node S only. Since $v_0 \in C$, therefore, v_0 cannot belong to C' and hence the edge-set of C' must intersect path P again at edge (v_0, v_1). Then, in that case, there are two edges (v_0, v_1) and (v_1, v_2) between $C \setminus C'$ and $C' \setminus C$. This implies that (C, C') are from Type-1, a contradiction. ◄

Generalizing Lemma 26 for any arbitrary length path, we give the following lemma which will provide the foundation for establishing 3-transversality of $(\lambda+1)$ (s,t)-cuts.

▶ **Lemma 27.** *Let $P = \langle v_0 \neq S, v_1, \ldots, v_i \rangle$ be a simple path in $\mathcal{D}_{\lambda+1}(G^U)$. There does not exist a $(\lambda+1)$ (s,t)-cut C such that $v_0, v_i \in C$ and $v_1, \ldots, v_{i-1} \in \overline{C}$.*

Figure 3 C: dash-dot curve, C': solid curve, $C \cup C'$: dashed curve, (v_{i-1}, v_i) is in $E(C) \cap E(C')$.

Proof. We begin with stating the following assertion.

$\mathcal{A}(i)$: There does not exist a $(\lambda+1)$ (s,t)-cut C such that $v_0, v_i \in C$ and $v_1, \ldots, v_{i-1} \in \overline{C}$ for a path $P = \langle v_0 \neq S, v_1, \ldots, v_{i-1}, v_i \rangle$ of length i.

We shall use proof by induction on the length of the path $i \geq 2$ to show that $\mathcal{A}(i)$ holds. The base case, $\mathcal{A}(2)$ follows directly from Lemma 26.

Let us now assume that $\mathcal{A}(j)$ holds for all $2 \leq j < i$. We shall prove that $\mathcal{A}(i)$ holds using a proof by contradiction. Assume to the contrary that there is a $(\lambda + 1)$ (s,t)-cut C such that $v_0, v_i \in C$ and $v_1, \ldots, v_{i-1} \in \overline{C}$ (refer to Figure 3). It follows from the construction of $\mathcal{D}_{\lambda+1}(G^U)$ that there is a $(\lambda + 1)$ (s,t)-cut C' for the edge (v_{i-1}, v_i) such that $v_{i-1} \in C'$ and $v_i \in \overline{C'}$. Along similar lines of the proof of Lemma 26 we can argue that (C, C') are pairs from Type-3 with $c(C \cup C') = \lambda + 1$, and edge-set of C' must be intersecting path P at least one more time before the edge (v_{i-1}, v_i). So, let (v_k, v_{k+1}), $0 \leq k < i - 1$, be the first edge on P that intersects edge-set of C'. Now, there are two possible cases, either $k = 0$ and $v_{k+1} = v_1$ or $k > 0$. In the former case, it is easy to observe that two edges (v_0, v_1) and (v_{i-1}, v_i) lie between $C' \setminus C$ and $C \setminus C'$. Therefore, (C, C') must be from Type-1, a contradiction. In the latter case, observe that $v_0, \ldots, v_k \notin C'$ and $v_{k+1} \in C'$ because (v_k, v_{k+1}) is the nearest edge from v_0 in P that intersects edge-set of C'. Therefore, $v_0, v_{k+1} \in C \cup C'$ and $v_1, \ldots, v_k \notin C \cup C'$. So, we get a path $P' = \langle v_0, \ldots, v_{k+1} \rangle$ and a $(\lambda+1)$ (s,t)-cut $C \cup C'$ such that $v_0, v_{k+1} \in C \cup C'$ and $v_1, \ldots, v_k \in \overline{C \cup C'}$. Moreover, path P' has length strictly smaller than i because $k < i - 1$. Therefore, $\mathcal{A}(k+1)$ fails to hold and $k + 1 < i$, a contradiction. ◀

Now in the following lemma, using Lemma 27 we argue that there cannot exist any $(\lambda+1)$ (s,t)-cut C in $\mathcal{D}_{\lambda+1}(G^U)$ such that edge-set of C intersects a path P more than thrice.

▶ **Lemma 28** (3-transversality property). *Each $(\lambda + 1)$ (s,t)-cut of G is 3-transversal in $\mathcal{D}_{\lambda+1}(G^I)$ or in $\mathcal{D}_{\lambda+1}(G^U)$.*

Proof. Without loss of generality let us consider $\mathcal{D}_{\lambda+1}(G^U)$. The proof is along similar lines for $\mathcal{D}_{\lambda+1}(G^I)$. We give a proof by contradiction. Assume to the contrary that there exists a $(\lambda+1)$ (s,t)-cut C in G and a path P in $\mathcal{D}_{\lambda+1}(G^U)$ such that edge-set of C intersects path P more than thrice. Then, path P can be divided into at least 5 contiguous disjoint subpaths $\{P_1, P_2, P_3, P_4, P_5\}$. For the cut C, observe that either each node of P_i is on side of S for all odd i and each node of P_i is on side of T for all even i or vice versa. In the former case, let us consider the subpath $P' = \langle P_3, P_4, P_5 \rangle$ of P. For this path P', we have each node of P_3, P_5 in C and each node of P_4 in $\overline{C}$. Moreover, S cannot belong to P_3 because by construction of $\mathcal{D}_{\lambda+1}(G^U)$ there is no incoming edge to S. Therefore, we have a path P' with first node (not S) and last node in $(\lambda+1)$ (s,t)-cut C and other nodes are in $\overline{C}$. This contradicts Lemma 27. In a similar way we can argue the latter case using subpath $P' = \langle P_2, P_3, P_4 \rangle$. ◀

■ **Figure 4** (i) The edge-set of $(\lambda + 1)$ (s,t)-cut A intersects path $\langle S, v_1, v_2, T \rangle$ thrice; A and B are $(\lambda + 1)$ (s,t)-cuts from Type-3. (ii) A 1-transversal cut $C = A \cap B$ with capacity greater than $\lambda + 1$.

We also show that there are $(\lambda + 1)$ (s,t)-cuts which may not appear in $\mathcal{D}_{\lambda+1}$ as 1-transversal cut (refer to Figure 4(i)). This is because, for each 3-transversal $(\lambda + 1)$ (s,t)-cut A, there exists a $(\lambda + 1)$ (s,t)-cut B such that A, B are pairs from Type-3. Therefore, our bound of 3-transversality is tight.

5.1.1   Extension to general graphs

The compact structure and characterization of $(\lambda + 1)$ (s,t)-cuts are valid if the graph has at most two (s,t)-mincuts. However, in general, graphs can have exponential number of (s,t)-mincuts. To tackle this difficulty we explore how a $(\lambda + 1)$ (s,t)-cut is related to the $(\lambda + 1)$ (s,t)-classes of a graph G. There are $(\lambda + 1)$ (s,t)-cuts in G which do not subdivide any $(\lambda + 1)$ (s,t)-class. We call them *degenerate* $(\lambda + 1)$ (s,t)-cuts. These cuts appear in $\mathcal{D}_\lambda$ and have the following characterization.

▶ **Lemma 29.** *An (s,t)-cut of G is a degenerate $(\lambda + 1)$ (s,t)-cut if and only if it appears as an (s,t)-cut with exactly one incoming edge in $\mathcal{D}_\lambda$.*

Henceforth, our main focus is on the non-degenerate $(\lambda + 1)$ (s,t)-cuts – cuts that subdivide $(\lambda + 1)$ (s,t)-classes. It follows from Lemma 12 that each such $(\lambda + 1)$ (s,t)-cut subdivides precisely one $(\lambda + 1)$ (s,t)-class. Therefore, the set of all $(\lambda + 1)$ (s,t)-cuts can be partitioned into disjoint subsets where each subset subdivides exactly one $(\lambda + 1)$ (s,t)-class. This partitioning allows us to work separately with each $(\lambda + 1)$ (s,t)-class. Let $\mathcal{W}$ be a $(\lambda + 1)$ (s,t)-class. In order to build a compact structure that stores all $(\lambda + 1)$ (s,t)-cuts that subdivide $\mathcal{W}$, we now define a graph $G(\mathcal{W})$ associated with $\mathcal{W}$ as follows.

Construction of $G(\mathcal{W})$:   Let τ be a topological ordering of $\mathcal{D}_\lambda$ where source node S (likewise T) has the smallest (likewise highest) topological number. $G(\mathcal{W})$ is obtained by forming a quotient graph of G using τ as follows. Let μ be the node of $\mathcal{D}_\lambda$ corresponding to $\mathcal{W}$. All nodes that precede μ in the topological ordering τ are contracted into a single source node S' and every node that succeeds μ are contracted to a single sink node T'.

It is easy to observe that $s \in S'$ (likewise $t \in T'$) since $s \in S$ (likewise $t \in T$). Henceforth without causing any ambiguity we denote an (S', T')-cut in $G(\mathcal{W})$ as an (s,t)-cut in $G(\mathcal{W})$. The following lemma establishes the mapping between the $(\lambda + 1)$ (s,t)-cuts of G and $G(\mathcal{W})$.

▶ **Lemma 30.** *A $(\lambda+1)$ (s,t)-cut C in G subdivides a $(\lambda+1)$ (s,t)-class $\mathcal{W}$ into $(\mathcal{W}_1, \mathcal{W} \setminus \mathcal{W}_1)$ if and only if there exists a $(\lambda + 1)$ (s,t)-cut $C' = \mathcal{W}_1 \cup \{S'\}$ in $G(\mathcal{W})$.*

It is easy to observe that graph $G(\mathcal{W})$ can have at most two (s,t)-mincuts – S' and complement of T'. Therefore, we construct pair of DAGs, $\mathcal{D}_{\lambda+1}((G(\mathcal{W}))^I)$ and $\mathcal{D}_{\lambda+1}((G(\mathcal{W}))^U)$, for each

$(\lambda+1)$ (s,t)-class $\mathcal{W}$. Now we summarize these structural and characterization results in the following theorem.

▶ **Theorem 31.** *For any directed multi-graph G, on n vertices and m edges with a designated source vertex s and a designated sink vertex t, there exists a 2-level DAG structure of $\mathcal{O}(m)$ size–(i) $\mathcal{D}_\lambda$ and (ii) a pair of DAGs associated with each node of $\mathcal{D}_\lambda$, such that all $(\lambda+1)$ (s,t)-cuts of G are compactly stored and characterized as follows.*

1. *a non-degenerate $(\lambda+1)$ (s,t)-cut of G subdividing a $(\lambda+1)$ (s,t)-class $\mathcal{W}$ is a 3-transversal cut in one of the two DAGs associated with $\mathcal{W}$, and*
2. *an (s,t)-cut of G is a degenerate $(\lambda+1)$ (s,t)-cut if and only if it has exactly one incoming edge in $\mathcal{D}_\lambda$.*

We now explore the possibility of answering query $Q(u,v)$ using the pair of DAGs, $\mathcal{D}_{\lambda+1}(G^I)$ or $\mathcal{D}_{\lambda+1}(G^U)$. Recall that for the case of (s,t)-mincuts, just storing a topological ordering of $\mathcal{D}_\lambda$ is sufficient to answer query $Q(u,v)$ because each 1-transversal cut in $\mathcal{D}_\lambda$ is also an (s,t)-mincut. However, a 1-transversal cut in $\mathcal{D}_{\lambda+1}(G^I)$ or $\mathcal{D}_{\lambda+1}(G^U)$ needs not be a $(\lambda+1)$ (s,t)-cut. For example, cut $A \cap B$ as shown in Figure 4(ii) is 1-transversal but has capacity $\lambda+2$. Note that A and B form a $(\lambda+1)$ (s,t)-cuts from Type-2. In the following section, we provide a compact data structure that can answer query $Q(u,v)$ even when (u,v) is not necessarily an edge.

6 Compact data structures for reporting $(\lambda+1)$ (s,t)-cuts

In this section we address the very fundamental problem of reporting any $(\lambda+1)$ (s,t)-cut C for a given pair of vertices $\{u,v\}$ such that $u \in C$ and $v \in \overline{C}$, if exists. In order to answer the query it is sufficient to verify if v is separated by at least one of the cuts from $N_1(u)$. Note that in case of (s,t)-mincut, it suffices to store the $N_0(u)$ which turns out to be unique for each vertex u. This structure occupies $\mathcal{O}(n^2)$ space and achieves $\mathcal{O}(|C|)$ time for reporting the (s,t)-mincut C such that $u \in C$ and $v \in \overline{C}$. Unfortunately, $N_1(u)$ can have more than one elements (as shown in Figure 4 that both A and B belong to $N_1(\nu)$). This is because unlike for (s,t)-mincuts, the set of all $(\lambda+1)$ (s,t)-cuts that keep a vertex u on the side of s is not closed under intersection and union operations.

In order to design compact data structure for reporting $(\lambda+1)$ (s,t)-cut, as discussed in Section 5.1.1, we work on each $(\lambda+1)$ (s,t)-class. Let $\mathcal{W}$ be a $(\lambda+1)$ (s,t)-class in graph G and $u,v \in \mathcal{W}$. Although $N_1(u)$ can have multiple elements, it can be shown using sub-modularity of cuts (Lemma 9) that $N_1(u,v)$ is unique. However, this fact alone guarantees $\mathcal{O}(|\mathcal{W}|^3)$ space data structure for determining the existence of $N_1(u,v)$ for any pair of vertices $u,v \in \mathcal{W}$. In order to achieve more compact data structure we explore how $N_1(u,v)$ is related to $N_1(u,w)$ for any $w \in \mathcal{W}$. The following lemma provides an important insight into this relationship.

▶ **Lemma 32.** *If $N_1(u,v) \neq N_1(u,w)$ for any $\{u,v,w\}$ in $\mathcal{W}$, then $\mathcal{W} \subseteq N_1(u,v) \cup N_1(u,w)$.*

Proof. The proof is by contradiction. Let $C = N_1(u,v)$ and $C' = N_1(u,w)$. Assume that there is a vertex in $\mathcal{W}$ which is also in $\overline{C \cup C'}$. In that case both $C \cap C'$ and $C \cup C'$ subdivide $\mathcal{W}$; therefore, capacity of each of them is strictly larger than λ. So it immediately follows from sub-modularity of cuts (Lemma 9) that $c(C \cap C')$ is $(\lambda+1)$. Therefore we have a $(\lambda+1)$ (s,t)-cut $C \cap C'$ which is a proper subset of at least one of C and C'. Therefore, C or C' fails to satisfy Definition 6 – a contradiction. ◀

Now let $N_1(u) = \{C_1, C_2, \ldots, C_l\}$.   Consider any vertex $v \in \mathcal{W}$. If v belongs to $\bigcap_{i=1}^{l} C_i$, then $N_1(u, v)$ does not exist; otherwise, Lemma 32 implies that v is separated from u by exactly one of the cuts from $N_1(u)$. As a result the sets $\overline{C_i} \cap \mathcal{W}$ for each $i \in [l]$ are disjoint. In order to determine the position of any other vertex, $w \in \mathcal{W}$, with respect to $N_1(u, v)$, we formulate the following query.

$$\text{BELONG}(u, v, w) = \begin{cases} 1 & \text{if } N_1(u, v) \text{ exists and } w \in N_1(u, v) \\ 0 & \text{otherwise.} \end{cases}$$

For each vertex $x \in \mathcal{W} \setminus \{u\}$, if $x \in \overline{C_i}$ for any $i \in [l]$ then mark x with label i. Now if label of v and w are same, then w is not present in $N_1(u, v)$, otherwise w belongs to $N_1(u, v)$. In this way we can answer query $\text{BELONG}(u, v, w)$ in $\mathcal{O}(1)$ time. Therefore, the following lemma is immediate.

▶ **Lemma 33.** *Let $\mathcal{W}$ be a $(\lambda + 1)$ (s, t)-class and $u \in \mathcal{W}$. There exists an $\mathcal{O}(|\mathcal{W}|)$ size data structure $\mathcal{N}_{\lambda+1}(u)$ that can determine in $\mathcal{O}(k)$ time whether there exists a $(\lambda + 1)$ (s, t)-cut C such that $u \in C$ and $v_1, \ldots, v_k \in \overline{C}$ for any $v_1, \ldots, v_k \in \mathcal{W}$. If C exists, the data structure can output $\overline{C} \cap \mathcal{W}$ in $\mathcal{O}(|\overline{C} \cap \mathcal{W}|)$ time.*

Moreover, given vertices $u, v_1, \ldots, v_k \in \mathcal{W}$, $\mathcal{N}_{\lambda+1}(u)$ can be used to report a $(\lambda + 1)$ (s, t)-cut C, if exists, in $\mathcal{O}(|\overline{C}|)$ time using an auxiliary $\mathcal{O}(n)$ space topological ordering of $\mathcal{D}_\lambda$ of G such that $u \in C$ and $v_1, \ldots, v_k \in \overline{C}$. Now based on Lemma 33 we construct the following data structure for $(\lambda + 1)$ (s, t)-class $\mathcal{W}$.

Description of $\mathcal{N}_{\lambda+1}$: It consists of $\mathcal{N}_{\lambda+1}(u)$ for each u in $\mathcal{W}$.

   $\mathcal{N}_{\lambda+1}$ for a $(\lambda + 1)$ (s, t)-class $\mathcal{W}$ occupies $\mathcal{O}(|\mathcal{W}|^2)$ space. Each $(\lambda + 1)$ (s, t)-class is disjoint from each other. Hence by constructing $\mathcal{N}_{\lambda+1}$ for each $(\lambda + 1)$ (s, t)-class of G we get an $\mathcal{O}(n^2)$ size data structure and it completes the proof of Theorem 3. We complement this result with a conditional lower bound of $\tilde{\Omega}(n^2)$ space based on Conjecture 4.

7   Dual edge sensitivity oracle for (s, t)-mincuts

In this Section we shall present an oracle that can efficiently report (s, t)-mincut upon the failure of a pair of edges in graph G. We say that an edge (u, v) belongs to a $(\lambda+1)$ (s, t)-class $\mathcal{W}$ if both $u, v \in \mathcal{W}$.

7.1   Handling dual edge failures

Consider the failure of two edges $e = (x, y), e' = (x', y')$ in G. Suppose at least one of $\{e, e'\}$ does not belong to any $(\lambda + 1)$ (s, t)-class. In this case the value of (s, t)-mincut decreases by at least 1 if $N_0(x, y)$ or $N_0(x', y')$ exists. The value of (s, t)-mincut decreases by exactly 2 if and only if both e, e' are contributing to a single (s, t)-mincut. To determine the existence of such an (s, t)-mincut, since (s, t)-mincuts are closed under union operation, it is sufficient to verify whether $y \notin N_0(x', y')$ and $y' \notin N_0(x, y)$. So we construct a data structure, denoted by $\mathcal{N}_\lambda$, which consists of $N_0(u)$ for each vertex u of G. This $\mathcal{O}(n^2)$ space data structure achieves $\mathcal{O}(1)$ time to report the resulting value of (s, t)-mincut on failure of $\{e, e'\}$ in this case.

   If e and e' belong to distinct $(\lambda + 1)$ (s, t)-classes, then it follows from Lemma 12 that the (s, t)-mincut value remains unchanged. Let us consider case when both e and e' belong to the same $(\lambda + 1)$ (s, t)-class, say $\mathcal{W}$. It follows as a simple corollary of Lemma 30 that we

Figure 5 Graph (i) has no $\lambda + 1$ (s, t)-cut containing edges (x, y) and (x', y'), but Graph (ii) does have one such cut shown by dashed curve.

just need to verify if there is a $(\lambda + 1)$ (s, t)-cut in $G(\mathcal{W})$ in which e and e' are contributing. Note that the only $(\lambda + 1)$ (s, t)-class of $G(\mathcal{W})$ is $\mathcal{W}$.

Consider the data structure $\mathcal{N}_{\lambda+1}$ for $\mathcal{W}$ in $G(\mathcal{W})$. For the existence of a $(\lambda + 1)$ (s, t)-cut in which both e and e' are contributing, observe that a necessary condition is that $y \notin N_1(x', y')$ and $y' \notin N_1(x, y)$. These conditions can be verified using $\mathcal{N}_{\lambda+1}$ in $\mathcal{O}(1)$ time. Upon checking these conditions, a natural approach would be to explore whether the union of these two cuts is a $(\lambda + 1)$ (s, t)-cut. Unfortunately, we can not infer anything conclusively from the union of these cuts as illustrated by the two graphs in Figure 5. In both these graphs, $N_1(x, y) \cup N_1(x', y')$ is not a $(\lambda + 1)$ (s, t)-cut . However, for graph (i), no $(\lambda + 1)$ (s, t)-cut exists that contains the two edges; for graph (ii), there is still a $(\lambda + 1)$ (s, t)-cut containing the two edges. Note that in these graphs, $N_1(x, y)$ and $N_1(x', y')$ are pairs of $(\lambda + 1)$ (s, t)-cuts from Type-2. Looking at this hurdle carefully, we get the following insight. Since y, y' both lie outside $N_1(x', y') \cup N_1(x, y)$, hence $c(N_1(x, y) \cup N_1(x', y'))$ has to be $> \lambda$. Therefore, if $c(N_1(x, y) \cap N_1(x', y'))$ is also $> \lambda$, then using sub-modularity of cuts (Lemma 9), their union is bound to be a $(\lambda + 1)$ (s, t)-cut and this will serve our purpose. Unfortunately, $G(\mathcal{W})$ does not ensure that $c(N_1(x, y) \cap N_1(x', y'))$ is greater than λ as shown in Figure 5.

In order to materialize the above insight, we use covering technique (Theorem 18) to build the pair of graphs $\{G(\mathcal{W})^I, G(\mathcal{W})^U\}$ that partition all $(\lambda + 1)$ (s, t)-cuts of $G(\mathcal{W})$. It follows from Theorem 18(2) that the capacity of the intersection (likewise union) of each pair of $(\lambda + 1)$ (s, t)-cut in $G(\mathcal{W})^I$ (likewise $G(\mathcal{W})^U$) is greater than λ. This is because the only (s, t)-mincut of $G(\mathcal{W})^I$ is $\overline{T'}$ and the only (s, t)-mincut of $G(\mathcal{W})^U$ is S'. Theorem 18(1) states that $\{G(\mathcal{W})^I, G(\mathcal{W})^U\}$ covers all the $(\lambda + 1)$ (s, t)-cuts of $G(\mathcal{W})$. Therefore, if $y' \notin N_1(x, y)$ and $y \notin N_1(x', y')$ in $G(\mathcal{W})^I$ or $x \notin N_1(y', x')$ and $x' \notin N_1(y, x))$ in $(G(\mathcal{W})^U)^{\mathcal{T}}$ then there is a $(\lambda + 1)$ (s, t)-cut in $G(\mathcal{W})$ in which e, e' are contributing. Now we shall establish the converse of this assertion.

Suppose there exists a $(\lambda + 1)$ (s, t)-cut C in $G(\mathcal{W})$ to which both edges (x, y) and (x', y') are contributing. Without loss of generality assume that C is present in $G(\mathcal{W})^I$. It can be observed that the cuts $N_1(x, y)$ and $N_1(x', y')$ in $G(\mathcal{W})^I$ are subsets of C. Hence $y \notin N_1(x', y')$ and $y' \notin N_1(x, y)$ in $G(\mathcal{W})^I$. If C is present in $G(\mathcal{W})^U$, exactly the same analysis can be carried out on $(G(\mathcal{W})^U)^{\mathcal{T}}$. So we can state the following lemma.

▶ **Lemma 34.** *A pair of edges $e = (x, y)$, $e' = (x', y')$ from $(\lambda + 1)$ class $\mathcal{W}$ are outgoing edges of a $(\lambda + 1)$ (s, t)-cut in $G(\mathcal{W})$ if and only if (i) $y \notin N_1(x', y')$ and $y' \notin N_1(x, y)$ in $G(\mathcal{W})^I$ or (ii) $x \notin N_1(y', x')$ and $x' \notin N_1(y, x)$ in $(G(\mathcal{W})^U)^{\mathcal{T}}$.*

Using the data structure $\mathcal{N}_{\lambda+1}$, it requires a constant number of BELONG queries to verify

the conditions mentioned in Lemma 34. Therefore, the data structure $\mathcal{F}$ for dual edge failure is as follows.

$\mathcal{F}$ consists of the following data structures:
- $\mathcal{N}_\lambda$ for graph G.
- $\mathcal{N}_{\lambda+1}$ for $G(\mathcal{W})^I$ and $(G(\mathcal{W})^U)^{\mathcal{T}}$ for each $(\lambda+1)$ (s,t)-class $\mathcal{W}$.

It can be observed that the resulting (s,t)-mincut value can be reported in $\mathcal{O}(1)$ time upon failure of any pair of edges from G using $\mathcal{F}$. Handling of dual edge insertion is covered in the full version of this paper. We summarize the results of this section in the following theorem.

▶ **Theorem 35.** *A directed multi-graph $G = (V, E)$, on $|V| = n$ vertices and $|E| = m$ edges with a designated source vertex s and a designated sink vertex t, can be preprocessed for constructing $\mathcal{O}(n^2)$ space Oracle $\{\mathcal{F}, \mathcal{I}\}$ that takes $\mathcal{O}(1)$ time to report the value of resultant (s,t)-mincut upon*

1. *failure of any given pair of edges $(x, y), (x', y') \in E$ using $\mathcal{F}$, or*
2. *insertion of any given pair of edges $(x, y), (x', y') \in V \times V$ using $\mathcal{I}$.*

▶ **Remark 36.** For a $(\lambda+1)$ (s,t)-class $\mathcal{W}$ and any pair of disjoint subsets $A, B \subset \mathcal{W}$, $\mathcal{F}$ can determine in $\mathcal{O}(|A||B|)$ time if there exists a $(\lambda+1)$ (s,t)-cut C such that $A \subseteq C$ and $B \subseteq \overline{C}$. If C exists, then it is possible to report C using $\mathcal{F}$ in $\mathcal{O}((|A|+|B|)|\mathcal{W}|+|\overline{C}|)$ time.

8   Conditional lower bound for dual edge sensitivity for (s, t)-mincuts

The problem of reachability in directed graph is as follows – Given a simple directed graph G with n vertices and m edges, preprocess it to form a data structure which can efficiently report if a given vertex v is reachable from another vertex u. The reachability in G is same as reachability in G_{SCC}, a directed acyclic graph which can be obtained by contracting each of the Strongly Connected Components to a single vertex. Henceforth, we shall assume that G is a directed acyclic graph. We transform the directed acyclic graph G into a graph $\mathcal{D}$ as follows.

Construction of $\mathcal{D}$: Create two additional vertices, namely s and t. Suppose Δ_v denotes the difference in the number of incoming and outgoing edges of any vertex v of G. For each vertex v in G, if $\Delta_v > 0$ we add Δ_v edges from v to t. Likewise, if $\Delta_v < 0$ we add Δ_v edges from s to v. Lastly, add two additional edge(s) from s to v and v to t for all v in G. Observe that the number of edges in this graph is only $\mathcal{O}(m)$. Thus, we state the following lemma.

▶ **Lemma 37.** *For a directed graph G with n vertices and m edges, there exists a directed acyclic multi-graph $\mathcal{D}$ with $\mathcal{O}(m)$ edges and $n + 2$ vertices such that a vertex v is reachable from a vertex u in G if and only if vertex v is reachable from vertex u in $\mathcal{D}$.*

The graph $\mathcal{D}$, that we have constructed, has a very interesting property. $\mathcal{D}$ is identical to the DAG $\mathcal{D}_\lambda$ that stores all (s,t)-mincuts in graph $\mathcal{D}$. We crucially exploit this property to derive an equivalence between reachability queries in graph $\mathcal{D}$ and dual edge failure (or insertion) query for (s,t)-mincut in $\mathcal{D}$. Since, the reachability structure of $\mathcal{D}$ and G is identical, we state the following lemma.

▶ **Lemma 38.** *Let G be a directed graph. A vertex v is reachable from a vertex u in G if and only if the value of (s,t)-mincut reduces by exactly 1 on removal of the edges $\{(s, u), (v, t)\}$ from graph $\mathcal{D}$ which is obtained from G using Lemma 37.*

Proof sketch. We know that $\mathcal{D}$ is same as $\mathcal{D}_\lambda$ of $\mathcal{D}$. So, removal of any edge from $\mathcal{D}$ reduces value of (s,t)-mincut by 1. Any (s,t)-cut in which both edges, $\{(s,u),(v,t)\}$, are contributing cannot be 1-transversal because of the u to v path. Every (s,t)-mincut is 1-transversal (Lemma 20), therefore, there cannot exist (s,t)-mincut in which both edges are contributing.

If there is no path from u to v, then both edges contribute to (s,t)-cut $C = \overline{R(\{u\}) \cup T}$; where $R(\{u\})$ defines the set of vertices reachable from u. We can show that C is also an (s,t)-mincut. Thus, upon removal of edges $\{(s,u),(v,t)\}$ the value of (s,t)-mincut reduces by 2. ◀

Using Conjecture 4 and Lemma 38 we state the following conditional lower bound.

▶ **Theorem 39.** *Assuming Directed Reachability Hypothesis holds, any data structure that can report the value of (s,t)-mincut for a designated source s and a designated sink t upon failure or addition of any pair of edges in a directed multi-graph with n vertices and m edges must either use $\tilde{\Omega}(n^2)$ space, or $\tilde{\Omega}(m)$ time.*

── **References** ──

1   Surender Baswana and Abhyuday Pandey. Sensitivity oracles for all-pairs mincuts. In Joseph (Seffi) Naor and Niv Buchbinder, editors, *Proceedings of the 2022 ACM-SIAM Symposium on Discrete Algorithms, SODA 2022, Virtual Conference / Alexandria, VA, USA, January 9 - 12, 2022*, pages 581–609. SIAM, 2022. `doi:10.1137/1.9781611977073.27`.

2   András A. Benczúr. A representation of cuts within 6/5 times the edge connectivity with applications. In *36th Annual Symposium on Foundations of Computer Science, Milwaukee, Wisconsin, USA, 23-25 October 1995*, pages 92–102. IEEE Computer Society, 1995. `doi:10.1109/SFCS.1995.492466`.

3   Keerti Choudhary. An optimal dual fault tolerant reachability oracle. In Ioannis Chatzigiannakis, Michael Mitzenmacher, Yuval Rabani, and Davide Sangiorgi, editors, *43rd International Colloquium on Automata, Languages, and Programming, ICALP 2016, July 11-15, 2016, Rome, Italy*, volume 55 of *LIPIcs*, pages 130:1–130:13. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2016. `doi:10.4230/LIPIcs.ICALP.2016.130`.

4   Camil Demetrescu, Mikkel Thorup, Rezaul Alam Chowdhury, and Vijaya Ramachandran. Oracles for distances avoiding a failed node or link. *SIAM J. Comput.*, 37(5):1299–1318, 2008. `doi:10.1137/S0097539705429847`.

5   Efim A Dinitz, Alexander V Karzanov, and Michael V Lomonosov. On the structure of the system of minimum edge cuts in a graph. *Issledovaniya po Diskretnoi Optimizatsii*, pages 290–306, 1976.

6   Yefim Dinitz. Maintaining the 4-edge-connected components of a graph on-line. In *Second Israel Symposium on Theory of Computing Systems, ISTCS 1993, Natanya, Israel, June 7-9, 1993, Proceedings*, pages 88–97. IEEE Computer Society, 1993. `doi:10.1109/ISTCS.1993.253480`.

7   Yefim Dinitz and Zeev Nutov. A 2-level cactus model for the system of minimum and minimum+1 edge-cuts in a graph and its incremental maintenance. In Frank Thomson Leighton and Allan Borodin, editors, *Proceedings of the Twenty-Seventh Annual ACM Symposium on Theory of Computing, 29 May-1 June 1995, Las Vegas, Nevada, USA*, pages 509–518. ACM, 1995. `doi:10.1145/225058.225268`.

8   Yefim Dinitz and Alek Vainshtein. The general structure of edge-connectivity of a vertex subset in a graph and its incremental maintenance. odd case. *SIAM J. Comput.*, 30(3):753–808, 2000. `doi:10.1137/S0097539797330045`.

9   Ran Duan and Seth Pettie. Dual-failure distance and connectivity oracles. In Claire Mathieu, editor, *Proceedings of the Twentieth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2009, New York, NY, USA, January 4-6, 2009*, pages 506–515. SIAM, 2009. URL: `http://dl.acm.org/citation.cfm?id=1496770.1496826`.

10   L. R. Ford and D. R. Fulkerson. Maximal flow through a network. *Canadian Journal of Mathematics*, 8:399–404, 1956. `doi:10.4153/CJM-1956-045-5`.

11   Harold N. Gabow. Efficient splitting off algorithms for graphs. In Frank Thomson Leighton and Michael T. Goodrich, editors, *Proceedings of the Twenty-Sixth Annual ACM Symposium on Theory of Computing, 23-25 May 1994, Montréal, Québec, Canada*, pages 696–705. ACM, 1994. `doi:10.1145/195058.195436`.

12   Zvi Galil and Giuseppe F. Italiano. Maintaining the 3-edge-connected components of a graph on-line. *SIAM J. Comput.*, 22(1):11–28, 1993. `doi:10.1137/0222002`.

13   Naveen Garg and Vijay V. Vazirani. A polyhedron with all s - t cuts as vertices, and adjacency of cuts. *Math. Program.*, 70:17–25, 1995. `doi:10.1007/BF01585926`.

14   Isaac Goldstein, Tsvi Kopelowitz, Moshe Lewenstein, and Ely Porat. Conditional lower bounds for space/time tradeoffs. In Faith Ellen, Antonina Kolokolova, and Jörg-Rüdiger Sack, editors, *Algorithms and Data Structures - 15th International Symposium, WADS 2017, St. John's, NL, Canada, July 31 - August 2, 2017, Proceedings*, volume 10389 of *Lecture Notes in Computer Science*, pages 421–436. Springer, 2017. `doi:10.1007/978-3-319-62127-2\_36`.

15   Ken-ichi Kawarabayashi and Mikkel Thorup. Deterministic edge connectivity in near-linear time. *J. ACM*, 66(1):4:1–4:50, 2019. `doi:10.1145/3274663`.

16   Yin Tat Lee and Aaron Sidford. Path finding methods for linear programming: Solving linear programs in õ(vrank) iterations and faster algorithms for maximum flow. In *55th IEEE Annual Symposium on Foundations of Computer Science, FOCS 2014, Philadelphia, PA, USA, October 18-21, 2014*, pages 424–433. IEEE Computer Society, 2014. `doi:10.1109/FOCS.2014.52`.

17   Thomas Lengauer and Robert Endre Tarjan. A fast algorithm for finding dominators in a flowgraph. *ACM Trans. Program. Lang. Syst.*, 1(1):121–141, 1979. `doi:10.1145/357062.357071`.

18   Jason Li. Deterministic mincut in almost-linear time. In Samir Khuller and Virginia Vassilevska Williams, editors, *STOC '21: 53rd Annual ACM SIGACT Symposium on Theory of Computing, Virtual Event, Italy, June 21-25, 2021*, pages 384–395. ACM, 2021. `doi:10.1145/3406325.3451114`.

19   Jason Li and Debmalya Panigrahi. Deterministic min-cut in poly-logarithmic max-flows. In Sandy Irani, editor, *61st IEEE Annual Symposium on Foundations of Computer Science, FOCS 2020, Durham, NC, USA, November 16-19, 2020*, pages 85–92. IEEE, 2020. `doi:10.1109/FOCS46700.2020.00017`.

20   Merav Parter. Dual failure resilient BFS structure. In Chryssis Georgiou and Paul G. Spirakis, editors, *Proceedings of the 2015 ACM Symposium on Principles of Distributed Computing, PODC 2015, Donostia-San Sebastián, Spain, July 21 - 23, 2015*, pages 481–490. ACM, 2015. `doi:10.1145/2767386.2767408`.

21   Merav Parter and David Peleg. Sparse fault-tolerant BFS structures. *ACM Trans. Algorithms*, 13(1):11:1–11:24, 2016. `doi:10.1145/2976741`.

22   Mihai Patrascu. Unifying the landscape of cell-probe lower bounds. *SIAM J. Comput.*, 40(3):827–847, 2011. `doi:10.1137/09075336X`.

23   Jean-Claude Picard and Maurice Queyranne. On the structure of all minimum cuts in a network and applications. *In Rayward-Smith V.J. (eds) Combinatorial Optimization II. Mathematical Programming Studies*, 13(1):8–16, 1980. `doi:10.1007/BFb0120902`.

24   Vijay V. Vazirani and Mihalis Yannakakis. Suboptimal cuts: Their enumeration, weight and number (extended abstract). In Werner Kuich, editor, *Automata, Languages and Programming, 19th International Colloquium, ICALP92, Vienna, Austria, July 13-17, 1992, Proceedings*, volume 623 of *Lecture Notes in Computer Science*, pages 366–377. Springer, 1992. `doi:10.1007/3-540-55719-9\_88`.